

RAHUL RACHHOYA

India | rahulrachhoya0@gmail.com | +91-8386067676
github.com/RahulRachhoya | linkedin.com/in/rahulrachhoya | rahulrachhoya.is-a.dev

AI Engineer with **3+ years** building production LLM systems (Claude, GPT-4, Mistral) serving 10K+ users. Specialized in **RAG, voice AI, multi-agent workflows** on AWS. **40% cost reduction, 35% hallucination improvement, 92% model accuracy**.

TECHNICAL SKILLS

LLMs & AI: Claude (Anthropic SDK), GPT-4, LLaMA, Gemini, Mistral, Fine-tuning (LoRA, PEFT), Prompt Engineering

Frameworks: LangChain, LangGraph, LlamaIndex, CrewAI, AutoGen, LangSmith, Streamlit, Hugging Face

RAG & DBs: FAISS, Pinecone, ChromaDB, Weaviate, pgvector, Qdrant, Hybrid Search, Cohere Reranking

Cloud & Ops: AWS (Bedrock, Lambda, EC2, S3, SageMaker), Docker, MLflow, Weights & Biases, Prometheus

Languages: Python, FastAPI, Flask, Django, REST APIs, pytest

PROFESSIONAL EXPERIENCE

AI Engineer - Voice AI & Conversational Systems

Jan 2026 - Present

Careers360

Hyderabad, India

- Architected voice AI career counselor using Claude/Bedrock + Sarvam AI (STT/TTS) + pgvector serving **10K+ sessions**; RAG pipeline reduces hallucinations **35%**
- Optimized costs **40% (\$0.40→\$0.24/session)** via prompt caching + A/B testing; LangSmith observability tracking latency, cost, quality
- Built multi-agent workflow (function calling) improving accuracy **28%**; AWS Lambda deployment with **99.7% uptime**

AI Engineer

Aug 2024 - Dec 2025

Crystaltech Services

Remote, India

- Built RAG pipelines reducing hallucinations **35%**; hybrid search (BM25+vector) + Cohere reranker improving retrieval **85%→91%**
- Fine-tuned **Mistral 7B** using LoRA achieving **92% accuracy** (vs 78% base) with **40% faster inference**; training cost **\$50 vs \$2K**
- Weights & Biases monitoring 50+ LLM iterations; A/B tested GPT-4 vs Claude achieving **60% cost savings**
- Built 4-agent document system (LangGraph) processing **500+ docs** with **94% accuracy**; **75% time reduction (30min→7min)**

System Engineer (AI/ML Focus)

Jun 2022 - Jul 2024

STL Digital Limited

Pune, India

- Implemented RAG system improving accuracy **40% (65%→91%)** using ChromaDB + GPT-3.5; 500+ daily queries with 1.8s response
- Built AI-assisted validation reducing errors **60%**; deployed Streamlit dashboard to AWS serving **500+ daily users**
- Implemented pytest suite with **85% coverage**; CI/CD pipeline via GitHub Actions

EDUCATION

Master of Computer Applications (MCA) - AI, Data Science, Machine Learning

2019 - 2022

University of Hyderabad, India | GPA: 8.2/10

CERTIFICATIONS & PUBLICATIONS

AWS ML Specialty (In Progress) | **Deep Learning Specialization** (Coursera)

Articles: "Building Production Voice AI with Claude" (5K+ views), "RAG at Scale" (3K+ views), "Reducing LLM Hallucinations" (2K+ views)

Open Source: LangChain (8 merged PRs), LlamaIndex (3 merged PRs)